

1. Let  $\Phi$  be a boolean formula in conjunctive normal form, with exactly three literals in each clause. Recall that an assignment of boolean values to the variables in  $\Phi$  **satisfies** a clause if at least one of its literals is TRUE. The **maximum satisfiability problem** for 3CNF formulas, usually called MAX3SAT, asks for the maximum number of clauses that can be simultaneously satisfied by a single assignment.

Solving MAX3SAT exactly is clearly also NP-hard; this question asks about approximation algorithms. Let  $\text{Max3Sat}(\Phi)$  denote the maximum number of clauses in  $\Phi$  that can be simultaneously satisfied by one variable assignment.

- (a) Suppose we assign variables in  $\Phi$  to be True or False using independent fair coin flips. Prove that the expected number of satisfied clauses is at least  $\frac{7}{8}\text{Max3Sat}(\Phi)$ .

**Solution:** Each clause contains three distinct literals. If two of those literals are a variable  $x$  and its negation  $\bar{x}$ , then the clause is satisfied by every assignment. Otherwise, the literals come from three distinct variables, whose values are chosen independently and uniformly at random. Thus, each clause of  $\Phi$  is satisfied with probability *at least*  $7/8$ . Linearity of expectation now implies that the expected number of satisfied clauses is *at least*  $7/8$  of the total number of clauses. The total number of clauses cannot be smaller than  $\text{Max3Sat}(\Phi)$ . ■

- (b) Let  $k^+$  denote the number of clauses satisfied by setting every variable in  $\Phi$  to TRUE, and let  $k^-$  denote the number of clauses satisfied by setting every variable in  $\Phi$  to FALSE. Prove that  $\max\{k^+, k^-\} \geq \text{Max3Sat}(\Phi)/2$ .

**Solution:** A clause is satisfied by setting every variable to TRUE if and only if it contains a positive literal, meaning a variable that is not negated. Thus,  $k^+$  is the number of clauses containing positive literals.

Similarly, a clause is satisfied by setting every variable to FALSE if and only if it contains a negative literal, meaning a negated variable. Thus,  $k^-$  is the number of clauses containing negative literals.

Every clause contains either a positive literal or a negative literal, so  $k^+ + k^-$  is greater than or equal to the total number of clauses. It follows that  $\max\{k^+, k^-\}$  is at least half the total number of clauses. The total number of clauses cannot be smaller than  $\text{Max3Sat}(\Phi)$ . ■

- (c) Let  $\text{Min3Unsat}(\Phi)$  denote the *minimum* number of clauses that can be simultaneously left *unsatisfied* by a single assignment. Prove that it is NP-hard to approximate  $\text{Min3Unsat}(\Phi)$  within a factor of  $10^{100}$ .

**Solution:** Suppose there is a polynomial-time  $10^{100}$ -approximation algorithm for MIN3UNSAT. Let  $\Phi$  be an arbitrary 3CNF formula. If  $\Phi$  is satisfiable, then  $\text{Min3Unsat}(\Phi) = 0$ , so the approximation algorithm must output 0. On the other hand, if  $\Phi$  is not satisfiable, then  $\text{Min3Unsat}(\Phi) > 0$ , so the output of the approximation algorithm must be at least  $\text{Min3Unsat}(\Phi)/10^{100} > 0$ . Thus,  $\Phi$  is satisfiable if and only if our approximation algorithm returns zero. In other words, we've just described a polynomial-time algorithm for 3SAT! ■


2. Consider the following algorithm for approximating the minimum vertex cover of a connected graph  $G$ : **Return the set of all non-leaf nodes of an arbitrary depth-first spanning tree.** (Recall that a depth-first spanning tree is a rooted tree; the root is not considered a leaf, even if it has only one neighbor in the tree.)

- (a) Prove that this algorithm returns a vertex cover of  $G$ .

**Solution:** Let  $T$  be a depth-first spanning tree of  $G$ , and let  $N$  be the set of nodes in  $G$  that are not leaves in  $T$ . A vertex  $v$  is a leaf in  $T$  if and only if the depth-first search visits every neighbor of  $v$  before it visits  $v$ . Thus, two leaves in  $T$  cannot be adjacent in  $G$ . Conversely, for any edge in  $G$ , at least one of its endpoints must be in  $N$ . We conclude that  $N$  is a vertex cover of  $G$ . ■

- (b) Prove that this algorithm returns a 2-approximation to the smallest vertex cover of  $G$ .

**Solution:** Let  $T$  be a depth-first spanning tree of  $G$ . We construct a *matching*  $M$  in  $T$ —a subgraph of  $T$  in which every node has degree 1—using the following recursive algorithm. Perform a preorder traversal of  $T$ , visiting the children of each node in arbitrary order; initially  $M$  is empty and every node is unmarked. Whenever the traversal visits an unmarked node  $v$  whose parent  $u$  is also unmarked, add edge  $uv$  to  $M$  and mark both  $u$  and  $v$ .<sup>1</sup>



The number of edges in the matching is at least half the number of non-leaf nodes in the tree.

By construction, every non-leaf node in  $T$  is marked and therefore adjacent to exactly one edge in  $M$ . (Some leaves are also marked, but that only helps us.) Thus, if  $T$  has  $k$  non-leaf nodes, then  $M$  contains at least  $k/2$  edges. Every vertex cover contains at least one node incident to each edge in  $M$ ; thus, the minimum vertex cover has at least  $k/2$  nodes. We conclude that our algorithm computes a vertex cover that is at most twice as large as the minimum vertex cover. ■

- (c) Describe an infinite family of connected graphs for which this algorithm returns a vertex cover of size *exactly*  $2 \cdot \text{OPT}$ . This family implies that the analysis in part (b) is tight. [Hint: First find just **one** such graph, with few vertices.]

**Solution:** Suppose the input graph  $G$  is a path of length  $2k$  with vertices  $v_0, v_1, \dots, v_{2k}$ . If we start our depth-first search at  $v_0$ , then  $v_{2k}$  is the only leaf in the depth-first spanning tree, so our algorithm returns a vertex cover of size  $2k$ . On the other hand, the odd-index vertices  $v_1, v_3, \dots, v_{2k-1}$  define a vertex cover of  $G$  of size  $k$ . ■

<sup>1</sup>This is just an implementation of the “dumb” approximation algorithm described in class, where we use DFS to discover unmarked edges.

3. Consider the following modification of the “dumb” 2-approximation algorithm for minimum vertex cover that we saw in class. The only change is that we return a set of edges instead of a set of vertices.

APPROXMINMAXMATCHING( $G$ ):

```

 $M \leftarrow \emptyset$ 
while  $G$  has at least one edge
     $uv \leftarrow$  any edge in  $G$ 
     $G \leftarrow G \setminus \{u, v\}$ 
     $M \leftarrow M \cup \{uv\}$ 
return  $M$ 

```

- (a) Prove that the output subgraph  $M$  is a *matching*—no pair of edges in  $M$  share a common vertex.

**Solution:** Let  $uv$  and  $vw$  be arbitrary edges in  $G$  that share a common vertex  $v$ . If the algorithm ever adds  $uv$  to  $M$ , it also immediately removes  $vw$  from  $G$ , so  $vw$  cannot be added to  $M$  later. Conversely, if the algorithm ever adds  $vw$  to  $M$ , it cannot add  $uv$  to  $M$  later. Thus, at most one of those two edges appears in  $M$ . ■

- (b) Prove that  $M$  is a *maximal* matching— $M$  is not a proper subgraph of another matching in  $G$ .

**Solution:** Every edge that is removed from  $G$  either belongs to  $M$  or shares a vertex with an edge in  $M$ . The algorithm terminates when all edges are removed. Thus, every edge that is not in  $M$  shares a vertex with at least one edge in  $M$ . ■

- (c) Prove that  $M$  contains at most twice as many edges as the *smallest* maximal matching in  $G$ .

**Solution:** For any two maximal matchings  $M$  and  $M'$ , each edge in  $M$  is incident to at most two edges in  $M'$ , and therefore  $|M| \leq 2|M'|$ . Let  $M'$  be the smallest maximal matching. ■

**Solution:** Every edge in  $M$  is incident to at least one matched vertex in any other maximal matching  $M'$ . Thus, the number of edges in  $M$  is at most the number of vertices in  $M'$ , which is twice the number of edges in  $M'$ . Let  $M'$  be the smallest maximal matching. ■

**Solution:** Let  $OPTMM$  denote the number of edges in the smallest maximal matching, let  $OPTVC$  denote the number of vertices in the smallest vertex cover.


The vertices of *any* maximal matching  $M'$  are a vertex cover for  $G$ . (If there is an uncovered edge, then  $M'$  is not maximal after all.) In particular, the vertices in the *smallest* maximal matching are a vertex cover for  $G$ . Thus,  $2 \cdot OPTMM \geq OPTVC$ .

We proved in class that the vertices of  $M$  are a 2-approximation of the smallest vertex cover. Thus,  $2 \cdot |M| \leq 2 \cdot OPTVC$ .

Combining these two inequalities gives us  $|M| \leq OPTVC \leq 2 \cdot OPTMM$ . ■

- (d) Describe an infinite family of graphs  $G$  such that the matching returned by APPROXMINMAXMATCHING( $G$ ) contains exactly twice as many edges as the smallest maximum matching in  $G$ . This family implies that the analysis in part (c) is tight. [Hint: First find just **one** such graph, with few vertices.]

**Solution:** Let  $G$  be the disjoint union of  $k$  paths of length 3; the smallest maximal matching in  $G$  consists of the middle edge of each path and therefore has size  $k$ . If  $\text{APPROXMINMAXMATCHING}(G)$  checks the first and last edge of each path before any of the middle edges. Then  $\text{APPROXMINMAXMATCHING}(G)$  returns a matching of size  $2k$  consisting of the first and last edge of each path.



The matching returned by  $\text{APPROXMINMAXMATCHING}(G)$  and the smallest maximal matching.

■